




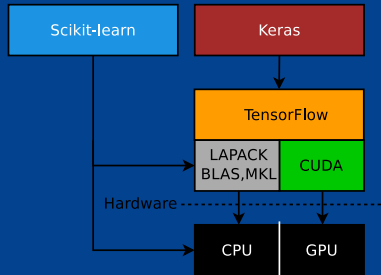
LESSON 7: Neural Networks II

CARSTEN EIE FRIGAARD

AUTUMN 2021



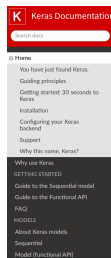

IMPLEMENTING MLPs IN KERAS



Keras and Tensorflow



Using the Keras API instead of Scikit-learn or TensorFlow



Docs » Home

Edit on GitHub

Keras: The Python Deep Learning library



You have just found Keras.

Keras is a high-level neural networks API, written in Python and capable of running on top of TensorFlow, CNTK, or Theano. It was developed with a focus on enabling fast experimentation. Being able to go from idea to result with the least possible delay is key to doing good research.

Use Keras if you need a deep learning library that:

- Allows for easy and fast prototyping (through user-friendliness, modularity, and extensibility).
- Supports both convolutional networks and recurrent networks, as well as combinations of the two.
- Runs seamlessly on CPU and GPU.

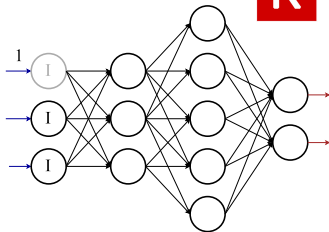
NOTE:

- ▶ documentation: <https://keras.io/>
- ▶ keras provides a `fit-predict`-interface,
- ▶ many similarities to Scikit-learn,
- ▶ but also many differences!

Building Keras MLPs



Using the Keras `Sequential` class,
programatically build up model:



```
1 # Build Keras model
2 model = Sequential()
3 model.add(Dense(input_dim=2, units=3, activation="tanh", ..)
4 model.add(Dense(units=5, activation="relu", ..)
5 model.add(Dense(units=2, activation="softmax"))
6 .
7 .
8 X_train, .. = train_test_split(X, y, .. )
9 .
10 .
11 y_train_categorical = to_categorical(y_train, num_classes=2)
12 y_test_categorical  = to_categorical(y_test,  num_classes=2)
13 .
14 .
15 history = model.fit(X_train, y_train_categorical, ..
16 .
17 .
18 score = model.evaluate(X_test, y_test_categorical)
```

Notes on Keras MLPs

Typical Keras MLP Supervised Classifier setup..

- ▶ loss function

```
loss='categorical_  
crossentropy'
```

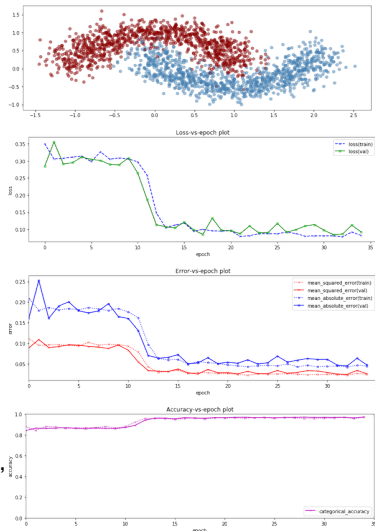
- ▶ metrics collected via history

```
metrics=[  
    'categorical_accuracy',  
    'mean_squared_error',  
    'mean_absolute_error'])
```

- ▶ input lay.: categorical encoding,

- ▶ output lay.: softmax function,

And notice that Keras do *not* provide metrics like
precision, recall, F1
but instead
categorical_accuracy, binary_accuracy



Input Layer: Categorical Encoding



For MLP Classification

One-hot to_categorical(.) encoding in Keras:

- ▶ input layer: one-hot class encoding,
- ▶ output layer: one neuron per output class that fires, and use `softmax` for output neurons
- ▶ beware of misformatted classes.

```
1 import numpy as np
2 from keras.utils.np_utils import to_categorical
3
4 y = np.array([1, 2, 0, 4, -1])
5 y_cat = to_categorical(y)
6
7 print(y_cat)
8
9 #[[0. 1. 0. 0. 0.] => i=0, class 1
10 # [0. 0. 1. 0. 0.] => i=1, class 2
11 # [1. 0. 0. 0. 0.] => i=2, class 0
12 # [0. 0. 0. 0. 1.] => i=3, class 4
13 # [0. 0. 0. 0. 1.] => i=4, also class 4!
14 # NOTE: no class 3
```

Output Layer: Softmax Function

For MLP Classification: Assigning a Probability for each Class

Softmax (softargmax/normalized exponential) definition

$$\text{softmax}(\mathbf{x})_i = \frac{e^{x_i}}{\sum_{i=1}^n e^{x_i}}$$

softmax: smooth approx. of **argmax** function.

argmax: the index-of-the-max-value for some data.

```
1 # python demo of softmax/argmax
2 x = np.array([1, 2, -4, 5, 1])
3 i = np.argmax(x)
```

```
4
5 PrintMatrix(x, "x = ")
6 print(f"np.argmax(x) = {np.argmax(x)}")
```

```
7
8 def softmax(x):
9     z = np.exp(x)
10    s = np.sum(z)
11    return z / s
```


```
12
13 PrintMatrix(softmax(x), "softmax(x) = ")
14 print(f"np.argmax(softmax(x)) = {np.argmax(softmax(x))}")
```

```
1 # output
2 x = [ 1  2 -4  5  1 ]
3 np.argmax(x) = 3
4 softmax(x) = [0.02 0.05 0. 0.92 0.02]
5 np.argmax(softmax(x)) = 3
```

KERAS INSTALL PARTY

Installing Keras

L07/Extras/installing_keras.ipynb



ITMAL Demo


Installing Keras

REVISIONS

2018-0325	CEF, initial.
2020-0305	CEF, F20 ITMAL update.
2020-0306	CEF, investigated Anaconda 2019.10 on Windows and updated GPU server notes.
2021-1012	CEF, updated for ITMAL E21.

Install Keras via Anaconda Prompt

1: Launch the **anaconda prompt** console (CMD), via the Start menu



High-Performance-Computing (HPC)

Running on the ASE GPU cluster..

- ▶ Your group login: `itmal09-e21`
- ▶ Password (can be changed): `itmal09-e21_123`

Documentation (always-out-of-date):

- ▶ Brightspace | ITMAL | Kursusinfo.. | GPU Cluster

[<https://brightspace.au.dk/d2l/le/lessons/27524/topics/296678>]



GPU CLUSTER CURRENTLY DOWN!

EXTRA SLIDES..

MLP Effect of Number of Hidden Layers

How Many Hidden Layers?

MNIST Search Quest Exercise:

- ▶ ITMAL Grp10:
 - used a 20-50-70-100-70-50-20 layer `MLPClassifier`,
- ▶ found layers by trial-and-error,
- ▶ but what are the optimal hidden layer sizes and neurons per hidden layer?




Figure 6.6: Empirical results showing that deeper networks generalize better when used to transcribe multi-digit numbers from photographs of addresses. Data from [Goodfellow et al. \(2014d\)](#). The test set accuracy consistently increases with increasing depth. See Fig. 6.7 for a control experiment demonstrating that other increases to the model size do not yield the same effect.

The ReLU Activation Function [DL]




Figure 6.3: The rectified linear activation function. This activation function is the default activation function recommended for use with most feedforward neural networks. Applying this function to the output of a linear transformation yields a nonlinear transformation. However, the function remains very close to linear, in the sense that it is a piecewise linear function with two linear pieces. Because rectified linear units are nearly linear, they preserve many of the properties that make linear models easy to optimize with gradient-based methods. They also preserve many of the properties that make linear models generalize well. A common principle throughout computer science is that we can build complicated systems from minimal components. Much as a Turing machine's memory needs only to be able to store 0 or 1 states, we can build a universal function approximator from rectified linear functions.

Effect of Depth [DL]




Figure 6.6: Empirical results showing that deeper networks generalize better when used to transcribe multi-digit numbers from photographs of addresses. Data from [Goodfellow *et al.* \(2014d\)](#). The test set accuracy consistently increases with increasing depth. See [Fig. 6.7](#) for a control experiment demonstrating that other increases to the model size do not yield the same effect.

Effect of Number of Parameters [DL]




Figure 6.7: Deeper models tend to perform better. This is not merely because the model is larger. This experiment from [Goodfellow *et al.* \(2014d\)](#) shows that increasing the number of parameters in layers of convolutional networks without increasing their depth is not nearly as effective at increasing test set performance. The legend indicates the depth of network used to make each curve and whether the curve represents variation in the size of the convolutional or the fully connected layers. We observe that shallow models in this context overfit at around 20 million parameters while deep ones can benefit from having over 60 million. This suggests that using a deep model expresses a useful preference over the space of functions the model can learn. Specifically, it expresses a belief that the function should consist of many simpler functions composed together. This could result either in learning a representation that is composed in turn of simpler representations (e.g., corners defined in terms of edges) or in learning a program with sequentially dependent steps (e.g., first locate a set of objects, then segment them from each other, then recognize them).

Deep feedforward networks [DL]

Deep feedforward networks, also often called feedforward neural networks, or multi-layer perceptrons (MLPs), are the quintessential deep learning models.

[DL, p167]