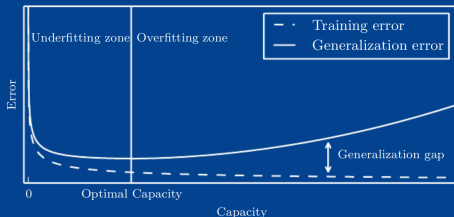




LESSON 9: Model-capacity, Under/over-fitting, Generalization

CARSTEN EIE FRIGAARD

AUTUMN 2021



"A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E ." — Mitchell (1997).

L09: Model-capacity, Under/over-fitting, Generalization

Agenda

- ▶ Résumé af GD og NN's.
- ▶ Model Capacity,
Exercise: [L09/capacity_under_overfitting.ipynb](#)
[OPTIONAL]
- ▶ Generalization Error,
Exercise: [L09/generalization_error.ipynb](#)
- ▶ Searching
Exercise: [L09/gridsearch.ipynb](#)

RESUMÉ: GD

The numerically Gradient decent [GD] method is based on the gradient vector

$$\nabla_{\mathbf{w}} J(\mathbf{w})$$

for the gradient oprator

$$\nabla_{\mathbf{w}} = \left[\frac{\partial}{\partial w_1}, \frac{\partial}{\partial w_2}, \dots, \frac{\partial}{\partial w_m} \right]^T$$


The algoritmn for updating via steps reads

$$\mathbf{w}^{(\text{next step})} = \mathbf{w} - \eta \nabla_{\mathbf{w}} J(\mathbf{w})$$


with η being the step size.

RESUMÉ: Training Deep Neural Networks

NN:



GD (via BPROP):



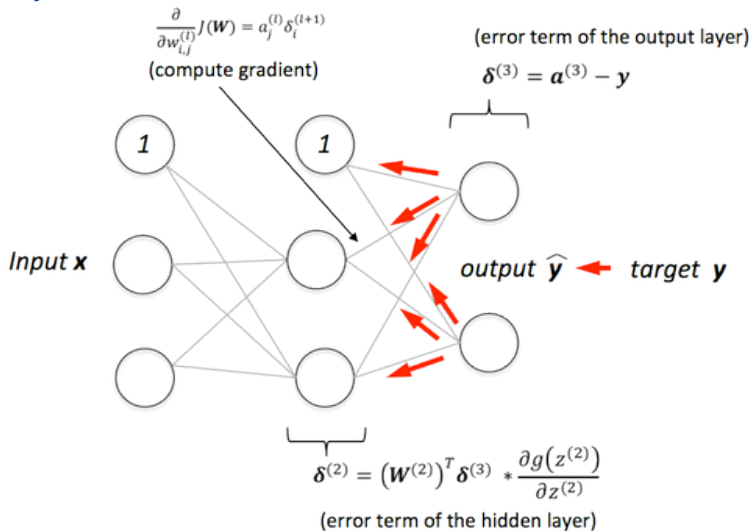
CONVEX \sim linear reg, not so for NNS

$$\mathbf{w}^{(\text{next})} = \mathbf{w} - \eta \nabla_{\mathbf{w}} J(\mathbf{w})$$

NOTE: NN: Neural net, GD: Gradient Descent, BPROP: Back Propagation

Backpropagation (BProp)

Training MLPs



NOTE: [<https://sebastianraschka.com/images/faq/visual-backpropagation/backpropagation.png>]

RESUMÉ: Training Deep Neural Networks

Equation 4-6. Gradient vector of the cost function

$$\nabla_{\theta} \text{MSE}(\theta) = \begin{pmatrix} \frac{\partial}{\partial \theta_0} \text{MSE}(\theta) \\ \frac{\partial}{\partial \theta_1} \text{MSE}(\theta) \\ \vdots \\ \frac{\partial}{\partial \theta_n} \text{MSE}(\theta) \end{pmatrix} = \frac{2}{m} \mathbf{X}^T (\mathbf{X}\theta - \mathbf{y})$$



Notice that this formula involves calculation set \mathbf{X} , at each Gradient Descent step! This is called *Batch Gradient Descent*: it uses the whole data at every step (actually, *Full Gradient Descent* would be a better name). As a result it is terribly slow on large sets (but we will see much faster Gradient Descent algorithms shortly). However, Gradient Descent scales well to many features; training a Linear Regression model on hundreds of thousands of features is much faster than using the Normal Equation or Stochastic Gradient Descent.

Once you have the gradient vector, which points uphill, you can move in the opposite direction to go downhill. This means subtracting $\nabla_{\theta} \text{MSE}(\theta)$ from the current parameters. The learning rate η comes into play:⁶ multiply the gradient vector by η to determine the size of the downhill step (Equation 4-7).

Equation 4-7. Gradient Descent step

$$\theta^{(\text{next step})} = \theta - \eta \nabla_{\theta} \text{MSE}(\theta)$$




Figure 11-1. Logistic activation function saturation





Figure 11-2. Leaky ReLU

$$\mathbf{w}^{(\text{next})} = \mathbf{w} - \eta \nabla_{\mathbf{w}} J(\mathbf{w})$$

MODEL CAPACITY



Model capacity

Exercise: `capacity_under_overfitting.ipynb`

Dummy and Paradox classifier:

capacity fixed ~ 0 , cannot generalize at all!

Linear regression for a polynomial model:

capacity \sim degree of the polynomial, x^n

Neural Network model:

capacity \propto number of neurons/layers

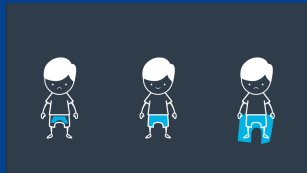
Homo sapiens ("modern humans"):

capacity \propto the IQ 'score' function?

\Rightarrow **Capacity** can be hard to express as a quantity for some models, but you need to choose..

\Rightarrow how to choose the **optimal capacity**?


UNDER- AND OVERFITTING



Under- and overfitting

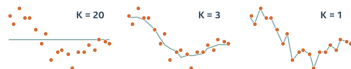
Exercise: `capacity_under_overfitting.ipynb`

Polynomial linear reg. fit for underlying model: $\cos(x)$



- ▶ underfitting:
capacity of model too low,
- ▶ overfitting:
capacity too high.

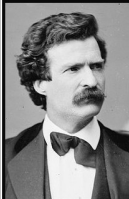
k-NN from L01:



⇒ how to choose the **optimal** capacity?

NOTE: HOML: Constraining a model [...] reduce risk of overfitting [via] regularization => L08

GENERALIZATION ERROR



All generalizations are false, including this one.

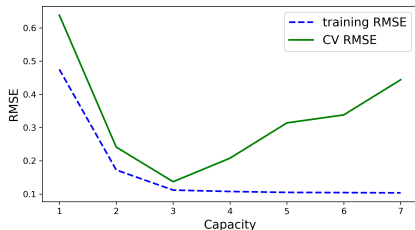
(Mark Twain)

Generalization Error


Exercise: `generalization_error.ipynb`

RMSE-capacity plot for lin. reg. with polynomial features

(capacity \sim degree of poly)



(Figure 5.3 from [DL])



Inspecting the plots from the exercise (`.ipynb`) and [DL],
extracting the concepts:

- ▶ training/generalization error,
- ▶ generalization gap,
- ▶ underfit/overfit zone,
- ▶ optimal capacity (best-model, early stop),
- ▶ (and the two axes: x/capacity, y/error.)

Generalization Error

Definition of ML:

“A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E .”

— Mitchell (1997).

Generalization Error

Exercise: generalization_error.ipynb

NOTE: three methods/plots:

- via **learning curves** as in [HOML],
- via an **error-capacity** plot as in [GITHOML] and [DL],
- via an **error-epoch** plot as in [GITHOML].

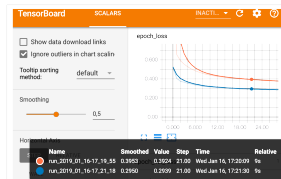





Figure 19.16. Visualizing Learning Curves with TensorBoard

SEARCHING

ML Algorithm +
Model Selection via Searching



ML Models (or ML algorithms)

Models encountered so far

Some classifiers and regressors..

```
sklearn.neighbors.KNeighborsRegressor  
sklearn.linear_model.LinearRegression  
sklearn.linear_model.SGDClassifier  
sklearn.linear_model.SGDRegressor
```

Perhaps..


```
sklearn.naive_bayes.GaussianNB  
sklearn.naive_bayes.MultinomialNB  
sklearn.svm.SVC  
sklearn.svm.SVR
```

and to some degree..

```
sklearn.linear_model.LogisticRegression  
sklearn.linear_model.Perceptron  
sklearn.neural_network.MLPClassifier  
sklearn.neural_network.MLPRegressor  
keras.Sequential
```

Or even more exotic models like..

- ▶ supervised ensemble: AdaBoost, Bagging, DecisionTree, RandomForest,...
- ▶ semi-supervised: ??
- ▶ unsupervised: K-means, manifolds, restricted Boltzmann machines,...
- ▶ clustering: K-means



ML Algorithm + Model Selection via Searching

What ML algorithm to choose?

- ▶ manual:
 - algorithm characteristics, \mathcal{O} complexity, etc.
 - browsing through Scikit-learn documentation,
 - ...and also based on data assumptions.
- ▶ semi-automatic:
 - brute-force model search, and fun with python!

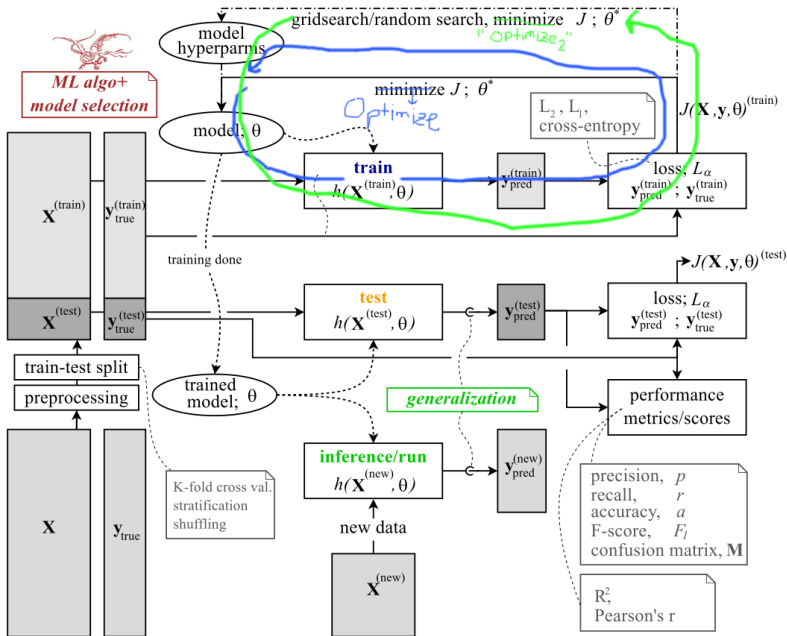
```
1 models = {  
2     SVC(gamma="scale"),  
3     SGDClassifier(tol=1e-3, eta0=0.1),  
4     GaussianNB()  
5 }  
6  
7 for i in models:  
8     i.fit(X_train, y_train)  
9     y_pred_test = i.predict(X_test)  
10    p = precision_score(y_test, y_pred_test, average='micro')  
11    print(f'{type(i).__name__:13s}: precision={p:0.2f}')
```

prints..

GaussianNB:	p=1.00
SGDClassifier:	p=0.93
SVC:	p=0.98

NOTE: Python set = {a, b}
Python dictionary = {a:x, b:y}

Model Selection via Grid Search



Model Selection via Grid Search

The hyperparameter-set for SGD linear regressor

```
1 class sklearn.linear_model.SGDRegressor(  
2     loss      = 'squared_loss', penalty      = 'l2',  
3     alpha     = 0.0001,               l1_ratio    = 0.15,  
4     tol       = None,                  shuffle     = True,  
5     verbose   = 0,                     epsilon     = 0.1,  
6     eta0      = 0.01,                  power_t     = 0.25,  
7     n_iter_no_change=5,                 warm_start  = False,  
8     fit_intercept = True,               max_iter    = None,  
9     average     = False,               n_iter      = None  
10    random_state  = None,               learning_rate='invscaling',  
11    early_stopping = False,             validation_fraction=0.1  
12 )
```



Search best hyperparameters in a (smaller) set, say


```
1 model = SGDRegressor()  
2 tuning_parameters = {  
3     'alpha': [ 0.001, 0.01, 0.1],  
4     'max_iter': [1, 10, 100, 1000],  
5     'learning_rate': ('constant', 'optimal', 'invscaling', 'adaptive')  
6 }  
7 ..  
8 grid_tuned = GridSearchCV(model, tuning_parameters, ..
```

Model Selection via Grid Search

How to select 'best' set of hyperparameter—using brute force?

Gridsearch seen in 3D for the two hyperspace dimensions:

- ▶ $\alpha \in [1, 2, 3, \dots]$ (NOTE: linear range for this plot only,
- ▶ $\text{max_iter} \in [1, 2, 3, \dots]$ should be 1, 10, 100 or similar.)




- ▶ why score and not J on z -axis?
- ▶ and what if there are many hyperparameters and many combinations? \rightarrow Zzzzzzz!

Model Selection via Randomized Search

How to select 'best' set of hyperparameters—faster than brute force?

Replace `GridSearchCV()` with


```
RandomizedSearchCV(n_iter=100,...)
```



- ▶ faster, but will not yield the (sub) optimal score maximum,
- ▶ ...but does it matter in a huge hyperparameter search-space?

Exercise: L09/gridsearch.ipynb

Qd MNIST Search Quest II: Husk at publicer på Brightspace

**Qd MNIST Search Quest II** ...

Publicer (helt nederst i denne item) løbende dit bedste resultat fra opgaven

gridsearch.ipynb (Qd)

Indtast din model constructor string og score og en kommentar vdr. dit model ala

```
Grp09: best: dat=mnist, score=0.90780, model=SGDClassifier(alpha=1.0,eta0=0.0001,learning_rate='invscaling')
```

```
Grp09: CTOR for best model: SGDClassifier(alpha=1.0, average=False, class_weight=None, early_stopping=False,
epsilon=0.1, eta0=0.0001, fit_intercept=True, l1_ratio=0.15,
learning_rate='invscaling', loss='hinge', max_iter=1000,
n_iter_no_change=5, n_jobs=None, penalty='l2', power_t=0.5,
random_state=None, shuffle=True, tol=0.001,
validation_fraction=0.1, verbose=0, warm_start=False)
```

NB: brug af **neurale netværks** modeller (Perceptrons, MLP's, Keras etc.) samt **KNeighborsClassifier** ikke tilladt i denne quest!

Highscore fra sidste semester = 0.97369

```
model=KNeighborsClassifier(algorithm='ball_tree', n_neighbors=3, p=4,weights='distance')
```

ITMAL Search Quest

Send message...

2/22